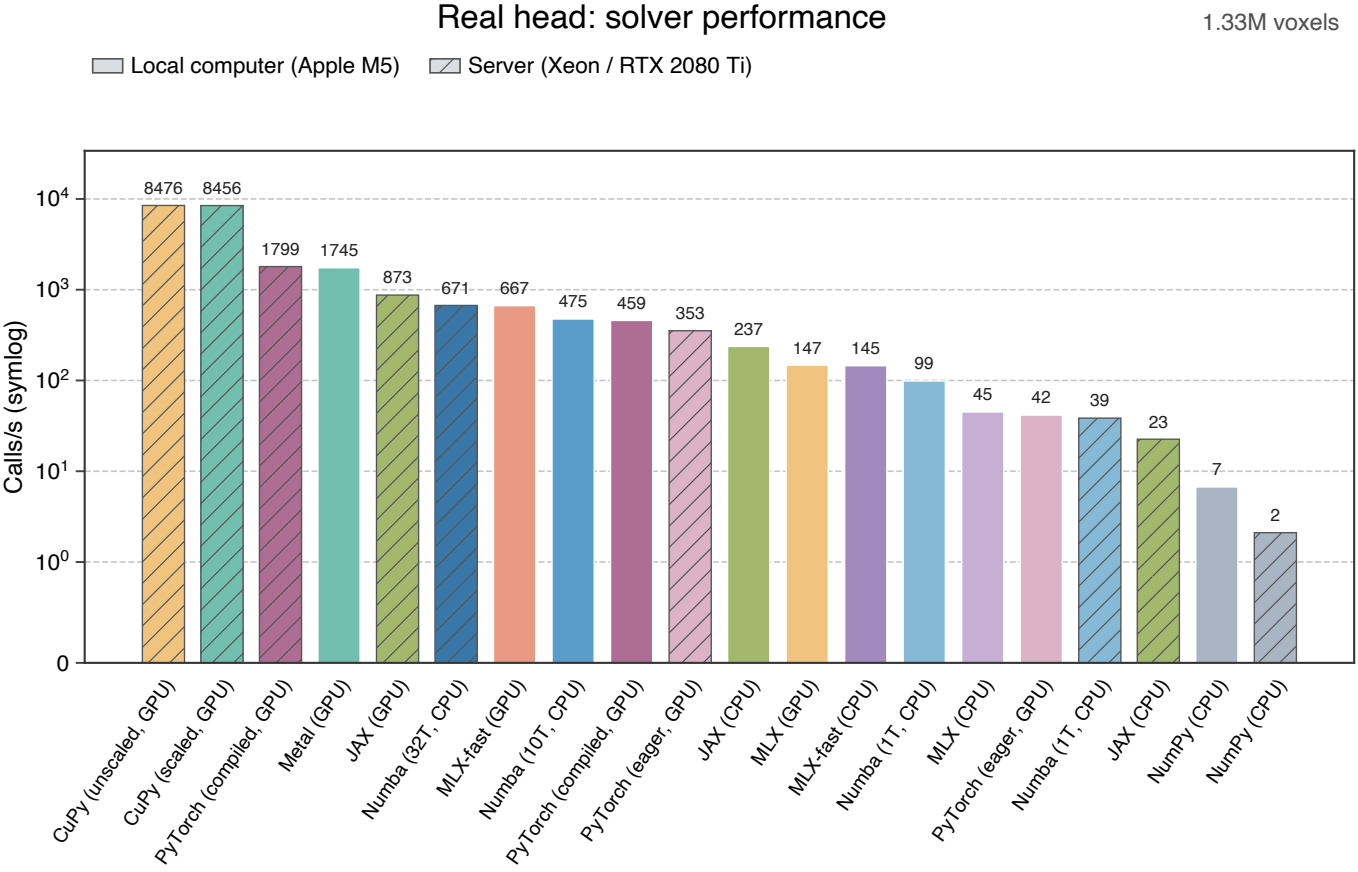
Real head: solver performance
1.33M voxels
Local computer (Apple M5)
Server (Xeon / RTX 2080 Ti)
Calls/s (symlog)
0
10^0
10^1
10^2
10^3
10^4
8476
8456
1799
1745
873
671
667
475
459
353
237
147
145
99
45
42
39
23
7
2
CuPy (unscaled, GPU)
CuPy (scaled, GPU)
PyTorch (compiled, GPU)
Metal (GPU)
JAX (GPU)
Numba (32T, CPU)
MLX-fast (GPU)
Numba (10T, CPU)
PyTorch (compiled, GPU)
PyTorch (eager, GPU)
JAX (CPU)
MLX (GPU)
MLX-fast (CPU)
Numba (1T, CPU)
MLX (CPU)
PyTorch (eager, GPU)
Numba (1T, CPU)
JAX (CPU)
NumPy (CPU)
NumPy (CPU)